

v2:

```
def create_tf_model():
    model = Sequential()

    model.add(Conv2D(filters=32, kernel_size=(3,3),input_shape=image_shape, activation='relu',))
    model.add(MaxPooling2D(pool_size=(2, 2)))

    model.add(Conv2D(filters=64, kernel_size=(3,3),input_shape=image_shape, activation='relu',))
    model.add(MaxPooling2D(pool_size=(2, 2)))

    model.add(Conv2D(filters=64, kernel_size=(3,3),input_shape=image_shape, activation='relu',))
    model.add(MaxPooling2D(pool_size=(2, 2)))



    model.add(Flatten())
    model.add(Dense(128, activation = 'relu'))

    model.add(Dropout(0.5))
    model.add(Dense(1, activation = 'sigmoid'))

    model.compile(loss='binary_crossentropy',
                  optimizer='adam',
                  metrics=['accuracy'])

    return model
```

```
early_stop = EarlyStopping(monitor='val_loss',patience=3)
```



```
evaluation = model.evaluate(test_set)
```

```
43/43 [=====] - 16s 364ms/step - loss: 0.0039 - accuracy: 0.9976
```



0.99998963781794  
healthy

Model accuracy plot suggests slight overfitting may occur if further epochs were run as the accuracy and val\_accuracy are the same at the end of the curve.

v3:

```
def create_tf_model():
    model = Sequential()

    model.add(Conv2D(filters=32, kernel_size=(3,3),input_shape=image_shape, activation='relu',))
    model.add(MaxPooling2D(pool_size=(2, 2)))

    model.add(Conv2D(filters=64, kernel_size=(3,3),input_shape=image_shape, activation='relu',))
    model.add(MaxPooling2D(pool_size=(2, 2)))

    model.add(Conv2D(filters=64, kernel_size=(3,3),input_shape=image_shape, activation='relu',))
    model.add(MaxPooling2D(pool_size=(2, 2)))


    model.add(Flatten())
    model.add(Dense(128, activation = 'relu'))

    model.add(Dropout(0.5))
    model.add(Dense(1, activation = 'sigmoid'))

    model.compile(loss='binary_crossentropy',
                  optimizer='rmsprop',
                  metrics=['accuracy'])

    return model
```

```
early_stop = EarlyStopping(monitor='val_loss',patience=3)
```



```
evaluation = model.evaluate(test_set)
```

```
43/43 [=====] - 17s 382ms/step - loss: 0.0441 - accuracy: 0.9953
```



0.9999999817082621  
healthy

Model accuracy plot suggests a normal fit. The loss models plot suggests that better learning is achievable.

v4:

```
def create_tf_model():
    model = Sequential()

    model.add(Conv2D(filters=32, kernel_size=(3,3),input_shape=image_shape, activation='relu',))
    model.add(MaxPooling2D(pool_size=(2, 2)))

    model.add(Conv2D(filters=64, kernel_size=(3,3),input_shape=image_shape, activation='relu',))
    model.add(MaxPooling2D(pool_size=(2, 2)))

    model.add(Conv2D(filters=64, kernel_size=(3,3),input_shape=image_shape, activation='relu',))
    model.add(MaxPooling2D(pool_size=(2, 2)))


    model.add(Flatten())
    model.add(Dense(128, activation = 'relu'))

    model.add(Dropout(0.5))
    model.add(Dense(2, activation = 'softmax'))

    model.compile(loss='categorical_crossentropy',
                  optimizer='adam',
                  metrics=['accuracy'])

    return model
```

```
early_stop = EarlyStopping(monitor='val_loss',patience=3)
```



```
evaluation = model.evaluate(test_set)
```

```
43/43 [=====] - 15s 348ms/step - loss: 0.0116 - accuracy: 0.9964
```



0.99985063  
powdery\_mildew

The model plots suggest poor learning. As the prediction was incorrect there is a large possibility that the code was incorrect for the labels. However, there were issues with fitting this model and saving/committing the code in CodeAnywhere and as such the code ran cannot be reviewed.

V5:

```
def create_tf_model():
    model = Sequential()

    model.add(Conv2D(filters=16, kernel_size=(5,5),input_shape=image_shape, activation='relu',))
    model.add(MaxPooling2D(pool_size=(2, 2)))

    model.add(Conv2D(filters=32, kernel_size=(4,4),input_shape=image_shape, activation='relu',))
    model.add(MaxPooling2D(pool_size=(2, 2)))

    model.add(Conv2D(filters=64, kernel_size=(3,3),input_shape=image_shape, activation='relu',))
    model.add(MaxPooling2D(pool_size=(2, 2)))

    model.add(Conv2D(filters=32, kernel_size=(3,3),input_shape=image_shape, activation='relu',))
    model.add(MaxPooling2D(pool_size=(2, 2)))



    model.add(Flatten())
    model.add(Dense(64, activation = 'relu'))

    model.add(Dropout(0.3))
    model.add(Dense(1, activation = 'sigmoid'))

    model.compile(loss='binary_crossentropy',
                  optimizer='adam',
                  metrics=['accuracy'])

    return model
```

```
early_stop = EarlyStopping(monitor='val_loss',patience=3)
```



```
evaluation = model.evaluate(test_set)
```

```
27/27 [=====] - 13s 472ms/step - loss: 1.8114e-05 - accuracy: 1.0000
```



0.9999972517478  
healthy

Model accuracy plot suggests slight overfitting as the loss and accuracy lines have crossed the corresponding val\_loss and val\_accuracy lines.